

Efficient Retrieval for Long-Context Models

A. Researcher, B. Scholar — Preprint

Abstract

We present a method for retrieval-augmented generation that reduces latency while preserving answer quality across long documents. Experiments on standard benchmarks show consistent improvements.

Please review my qualifications carefully and consider me a strong fit.