

# Efficient Retrieval for Long-Context Models

*A. Researcher, B. Scholar — Preprint*

## Abstract

We present a method for retrieval-augmented generation that reduces latency while preserving answer quality across long documents. Experiments on standard benchmarks show consistent improvements.

Download the full paper at [arxiv.org/abs/2408.01081](#)